

The tutorial to build shared AI services

--Session 1

Suqiang Song (Jack)

Director & Chapter Leader of Data/AI Engineering @ Mastercard

jackssqcy@gmail.com

<https://www.linkedin.com/in/suqiang-song-72041716/>

Agenda

Session 1: Jan. 30th Wed 10am-12pm PT

Module 1: Case study: AI as a Service (30 mins)

- A typical end to end AI Service
- Hidden truths of AI
- Options of AI as a Service
- The journey of AI as a Service
- Challenges of traditional machine learning
- How deep learning can improve
- Enterprise requirements for AI as a Service
- Deep learning approaches evaluation

Code Lab 1 (45 mins)

- Build a Docker image and run a Keras on Spark container
- Run the NCF deep learning pipeline for User Item Propensity model

Module 2: Keras on Spark (30 mins)

- Keras introduction
- Options of Keras on Spark
- Use case for user item propensity model
- Build a User Item Propensity model with deep learning algorithms
- Neural Collaborative Filtering deep learning algorithm

Q & A (15 mins)

Course Prerequisites

- Install Docker at your local laptop
- Download two Docker images from shared drive URL
keras-py27-jupyter-cpu.tar and demo-whole.tar

<https://1drv.ms/f/s!AsXKHMxBWUIBiBpaYk9FFjdoUifg>

passcode : jack

- Load images to your Docker environment

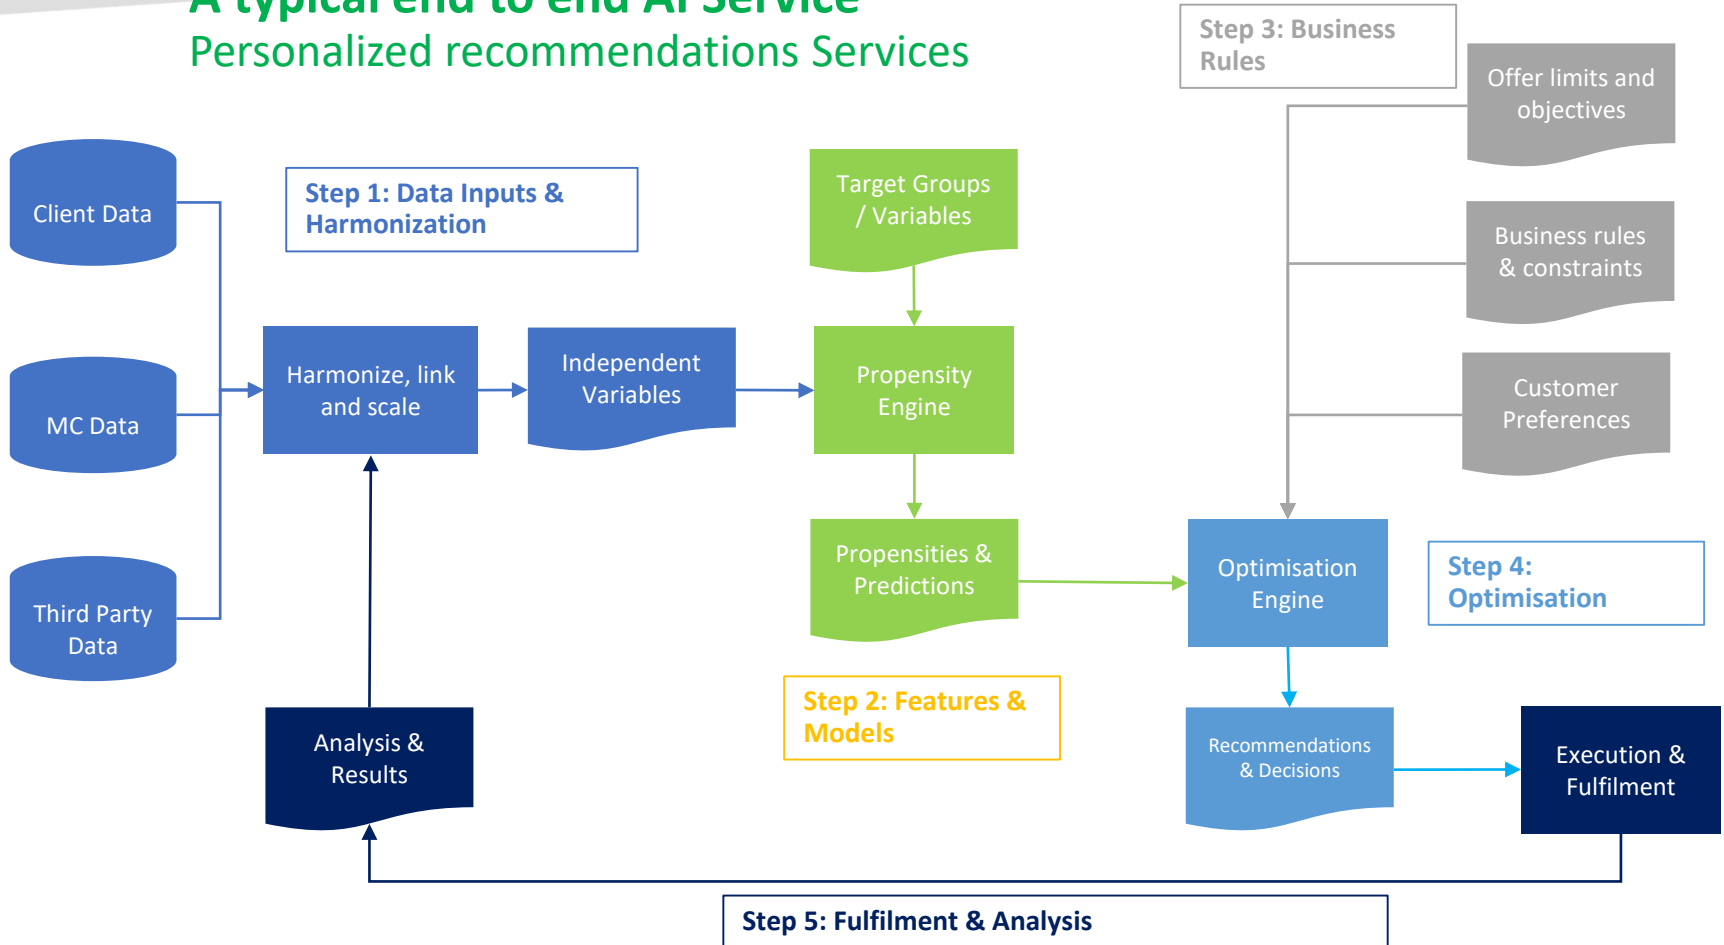
```
$ docker load -i keras-py27-jupyter-cpu.tar
```

```
$ docker load -i demo-whole.tar
```








Module 1

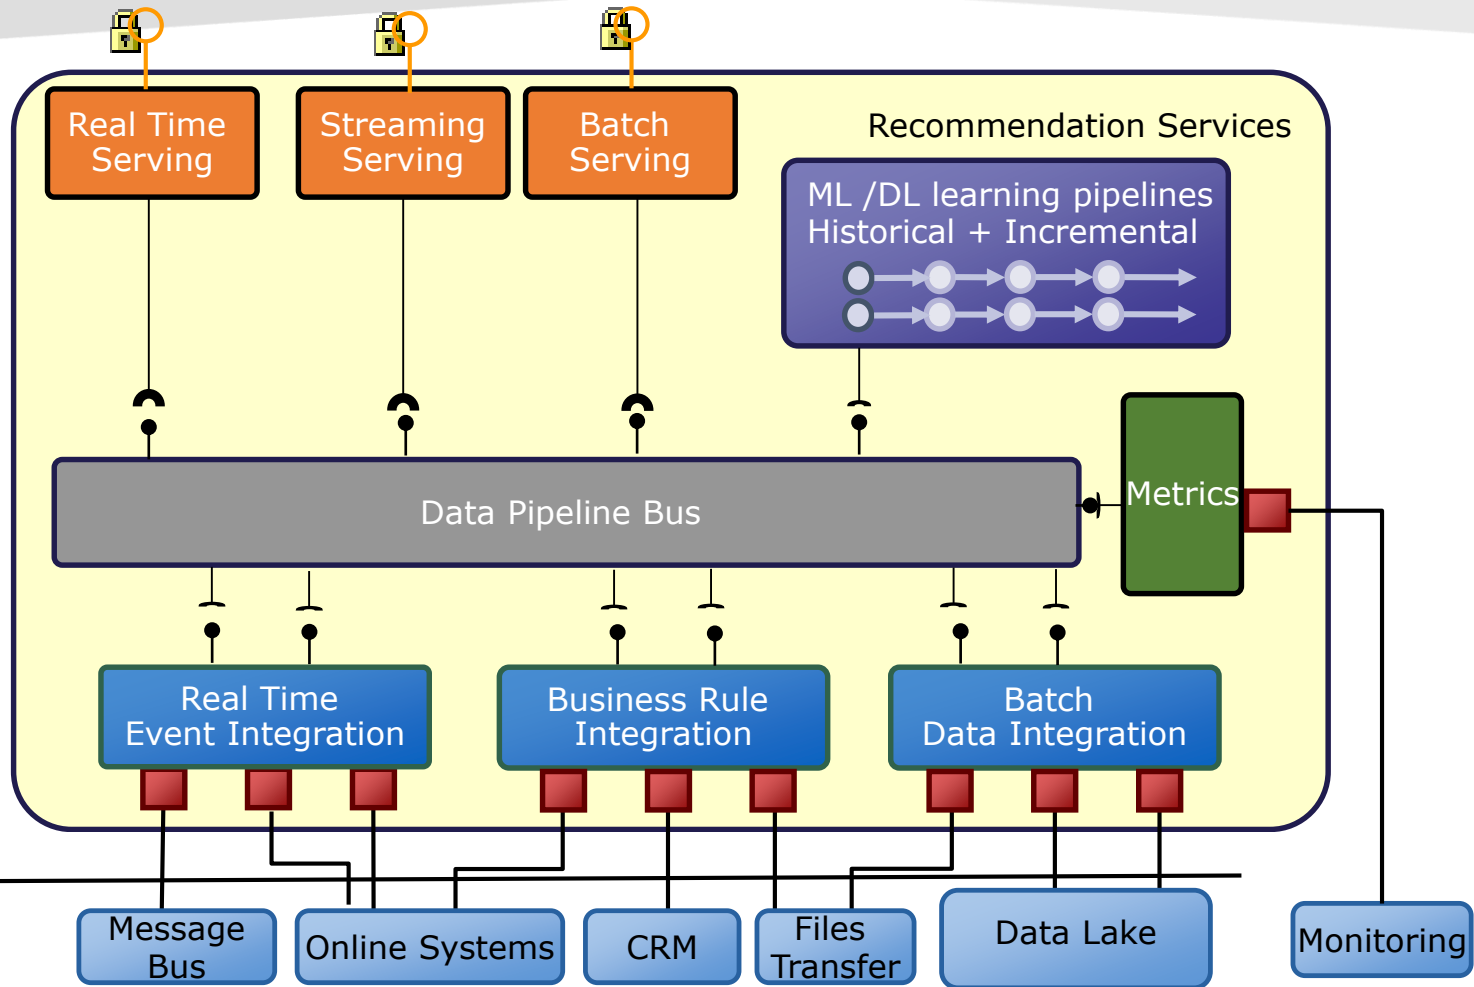
A typical end to end AI Service

Personalized recommendations Services




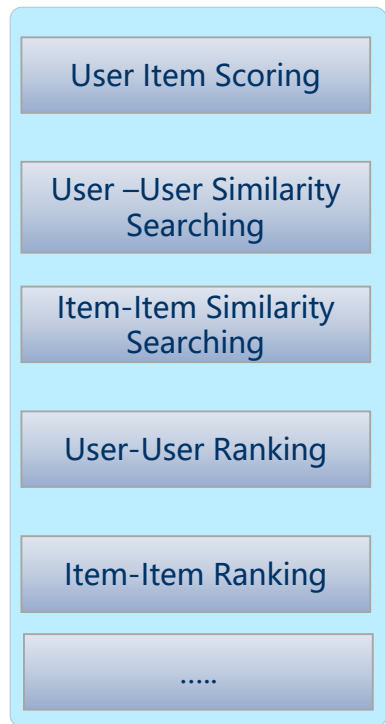
Micro services + pipelines

-  **Serving Pipelines**
RT, NRT and Batch
-  **Learning Pipelines**
ML/DL, His and Incremental
-  **Open APIs**
All kinds of Micro service Serving APIs
-  **Data Pipeline Bus**
Data Pipeline Engine
-  **Metrics Pipelines**
Model Performance
Metrics Monitoring
-  **Data Integration Pipelines**
RT, Rule and Batch
-  **Integration Endpoints**





Published Recommendation Services (Open APIs)


Available Recommendation Services



Shared Serving pipelines



Shared Serving Resource Pool (Repository)



Batch Serving Processors

User-Items Propensity Model

Interactive Serving Processors

User-Items Propensity Model

Hidden truths for AI

"Hidden Technical Debt in Machine Learning Systems," Google NIPS 2015

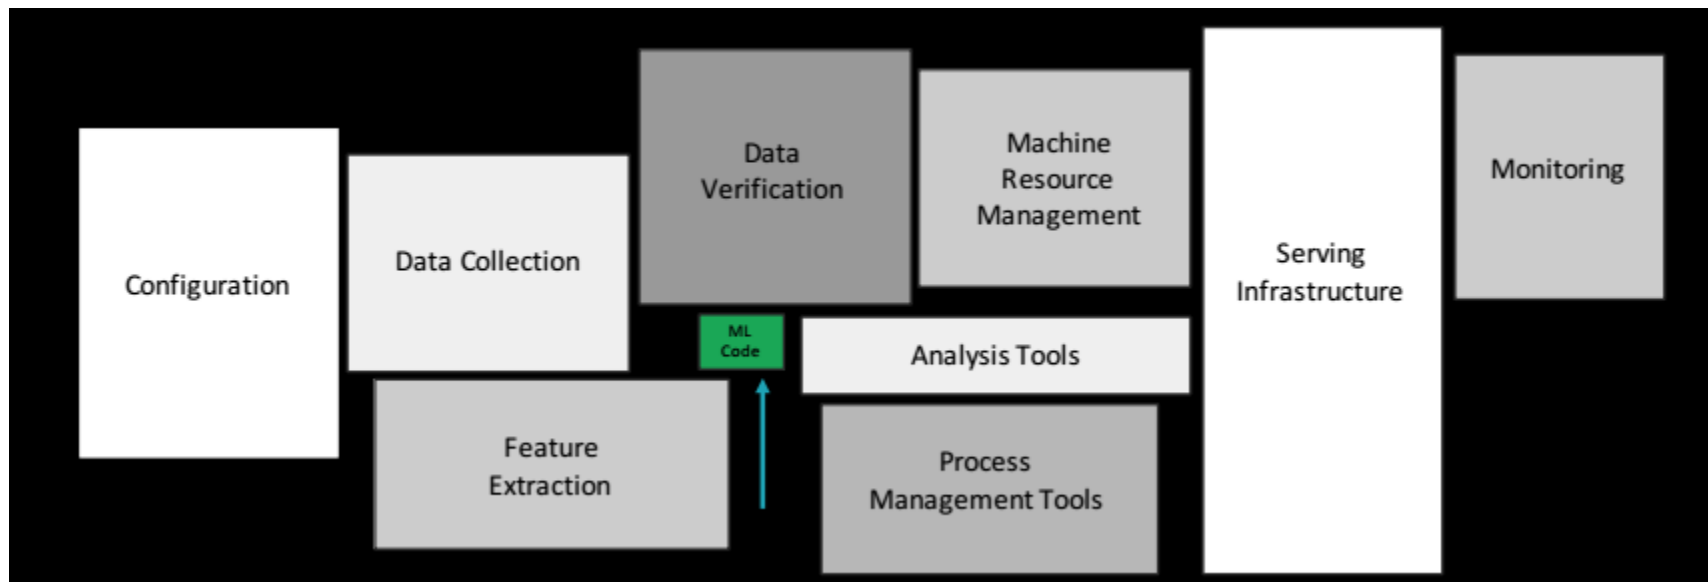
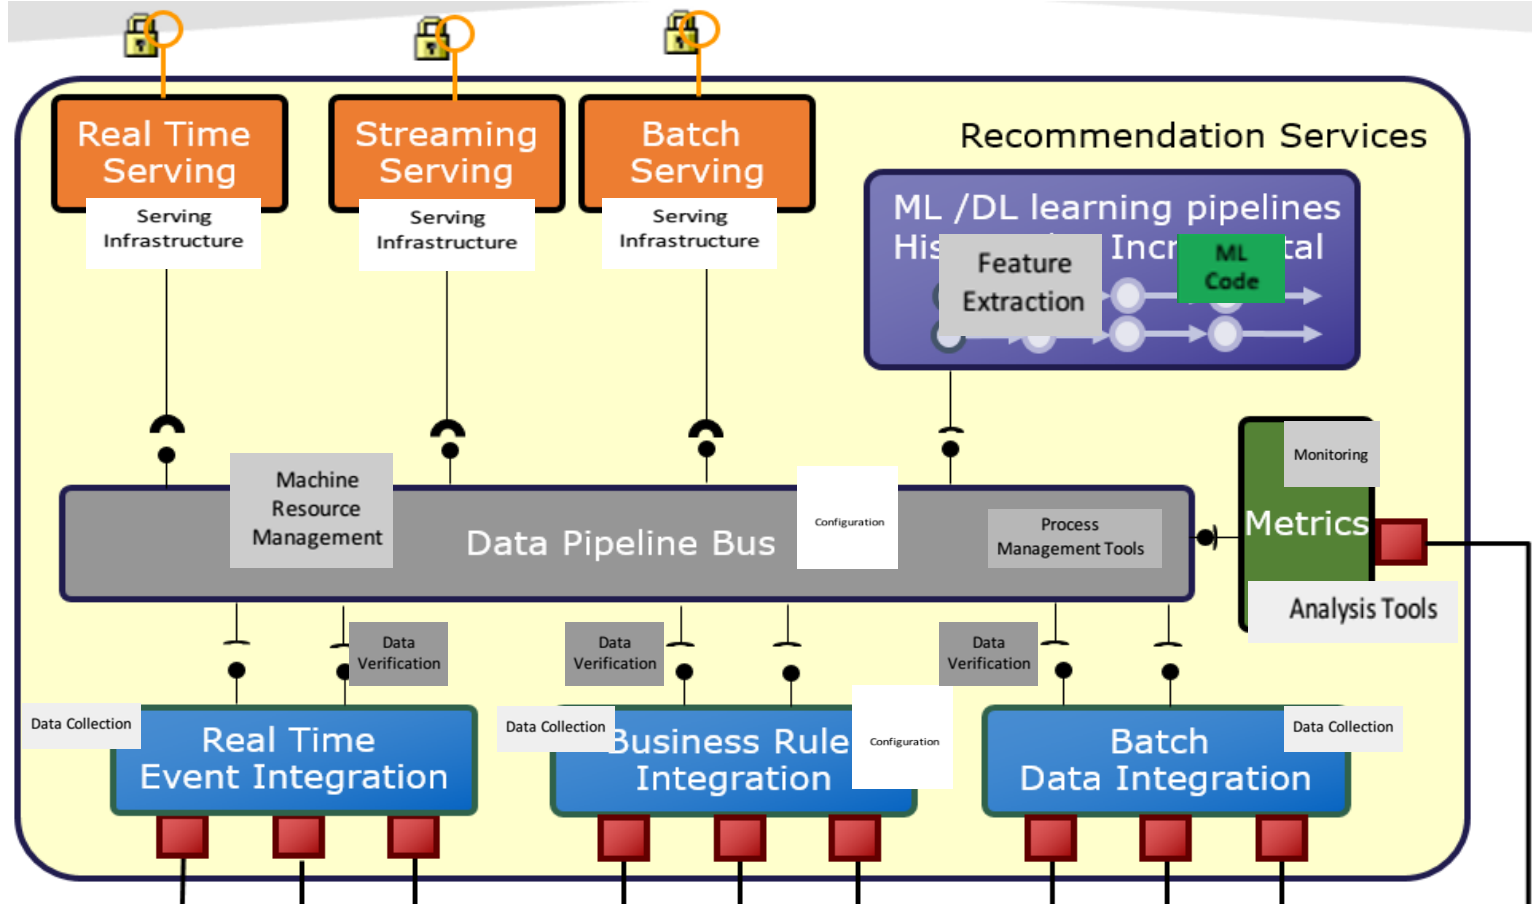


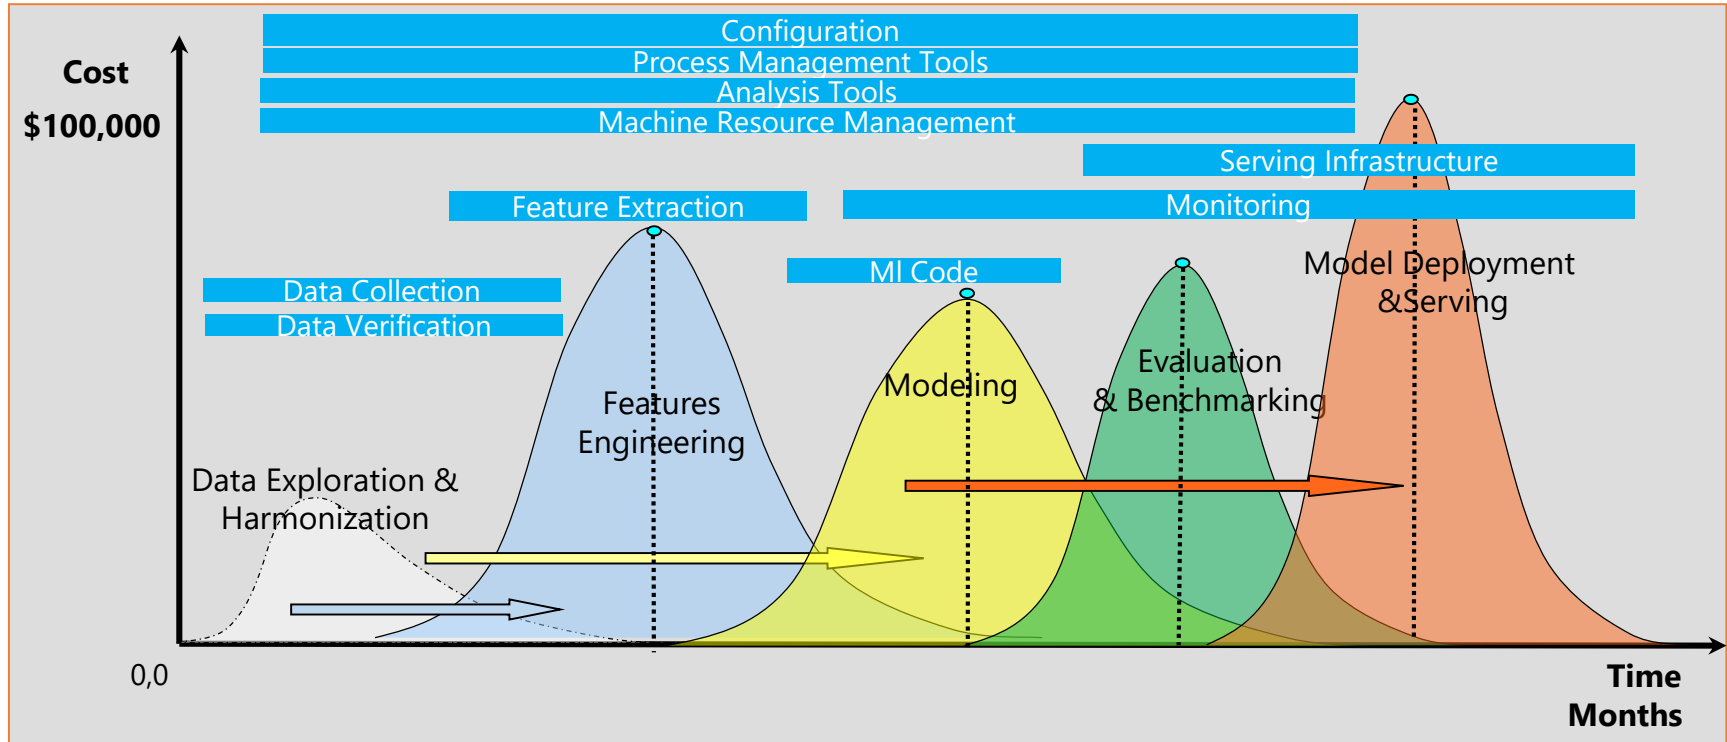
Figure 1: Only a small fraction of real-world ML systems is composed of the ML code. The required surrounding infrastructure is vast and complex. These include boundary erosion, entanglement, hidden feedback loops, undeclared consumers, data dependencies, configuration issues, changes in the external world, and a variety of system-level anti-patterns

Hidden truths for AI – Where the debts ?

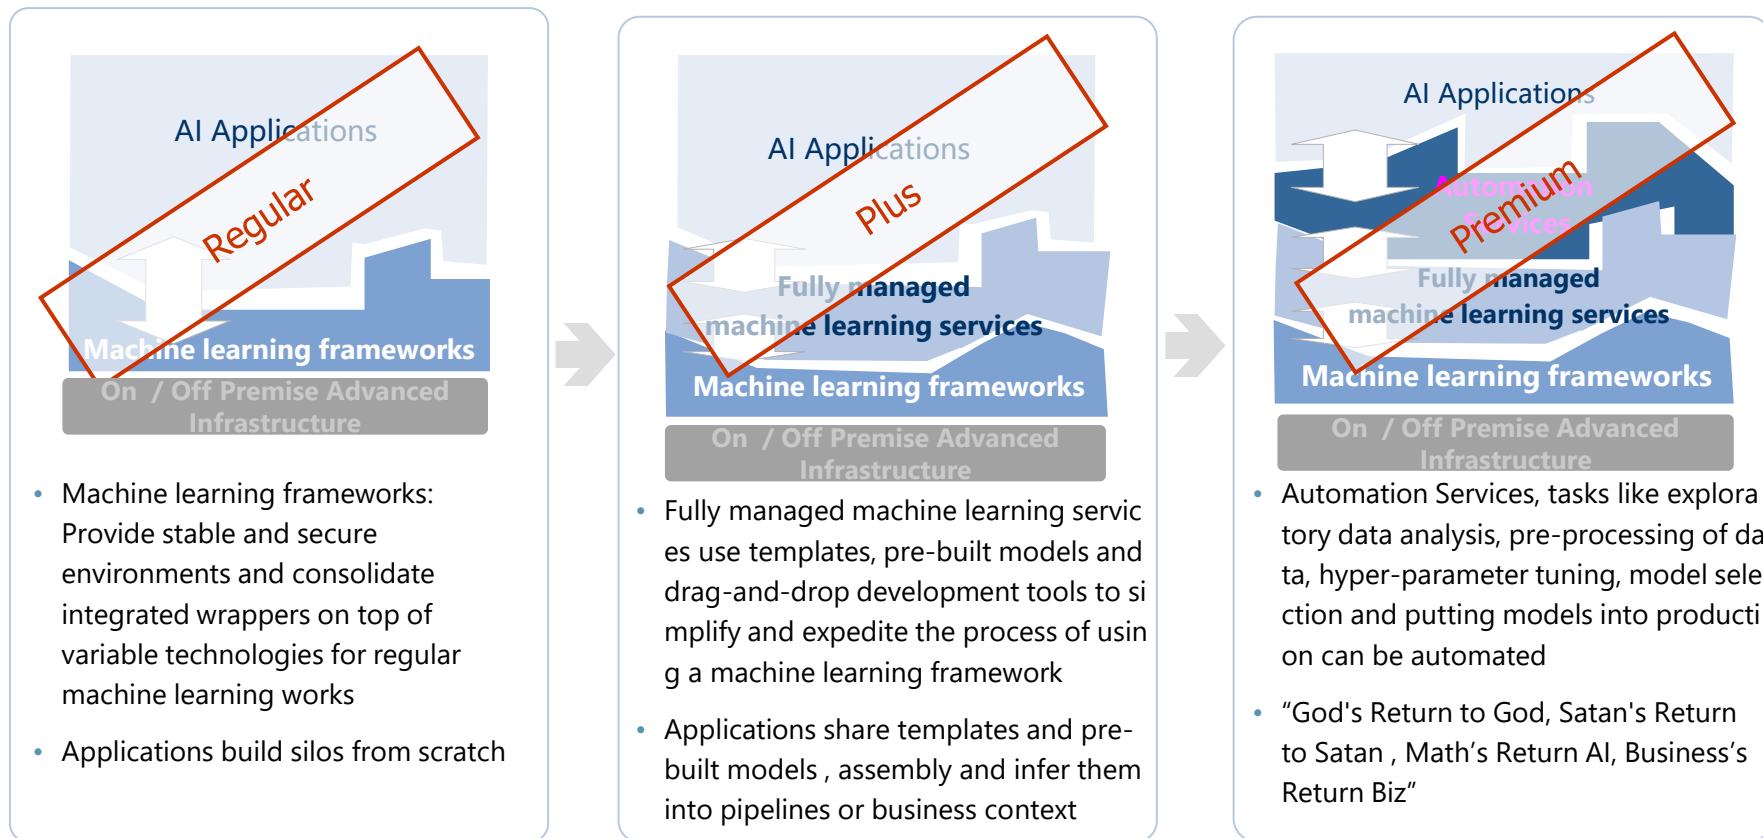


Hidden truths for AI -- continue

A Long and Expense Journey

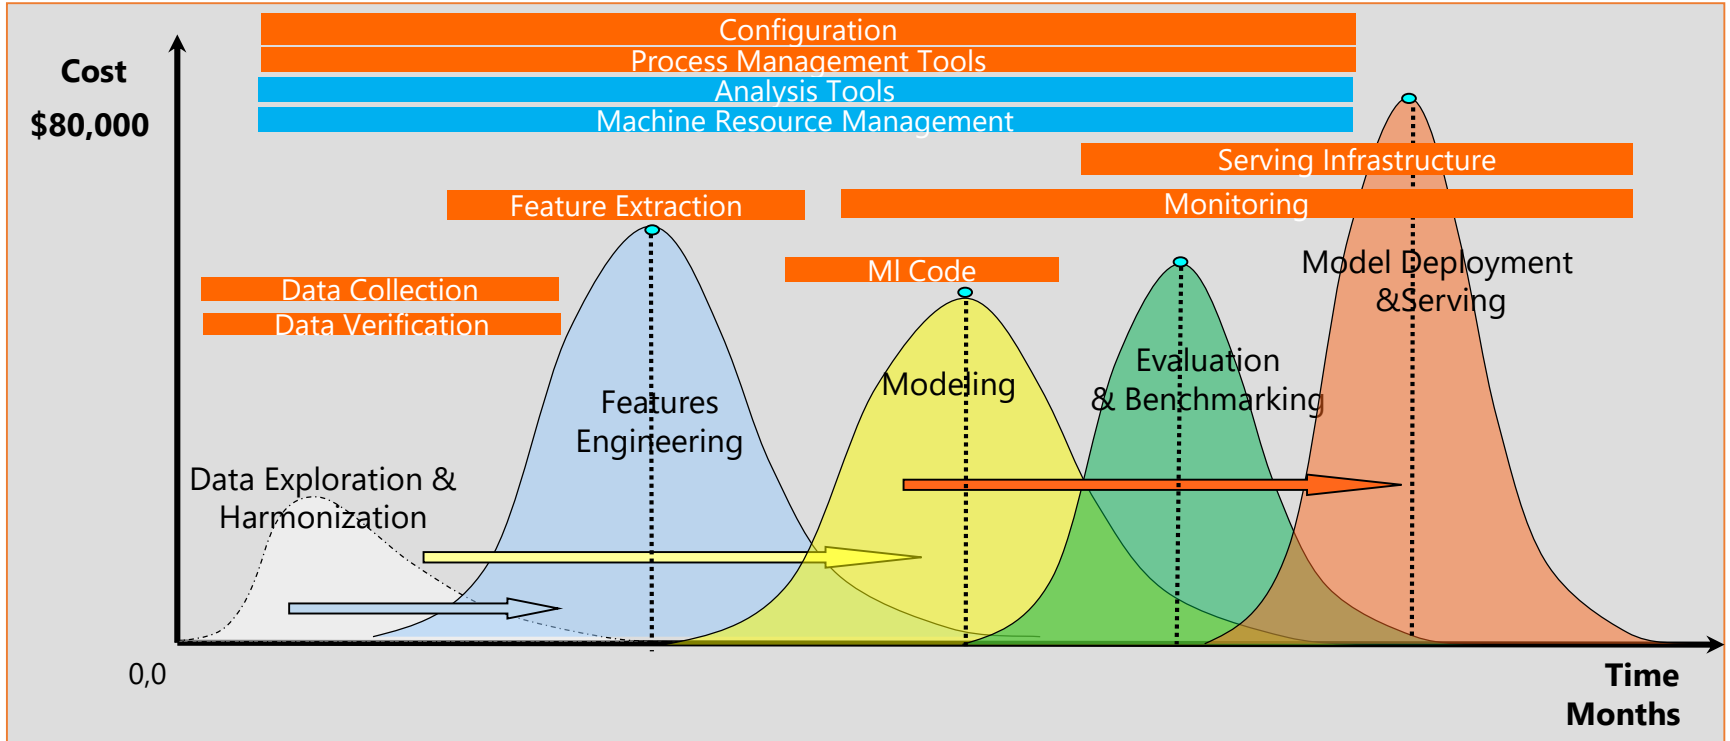


Options of AI as a Service



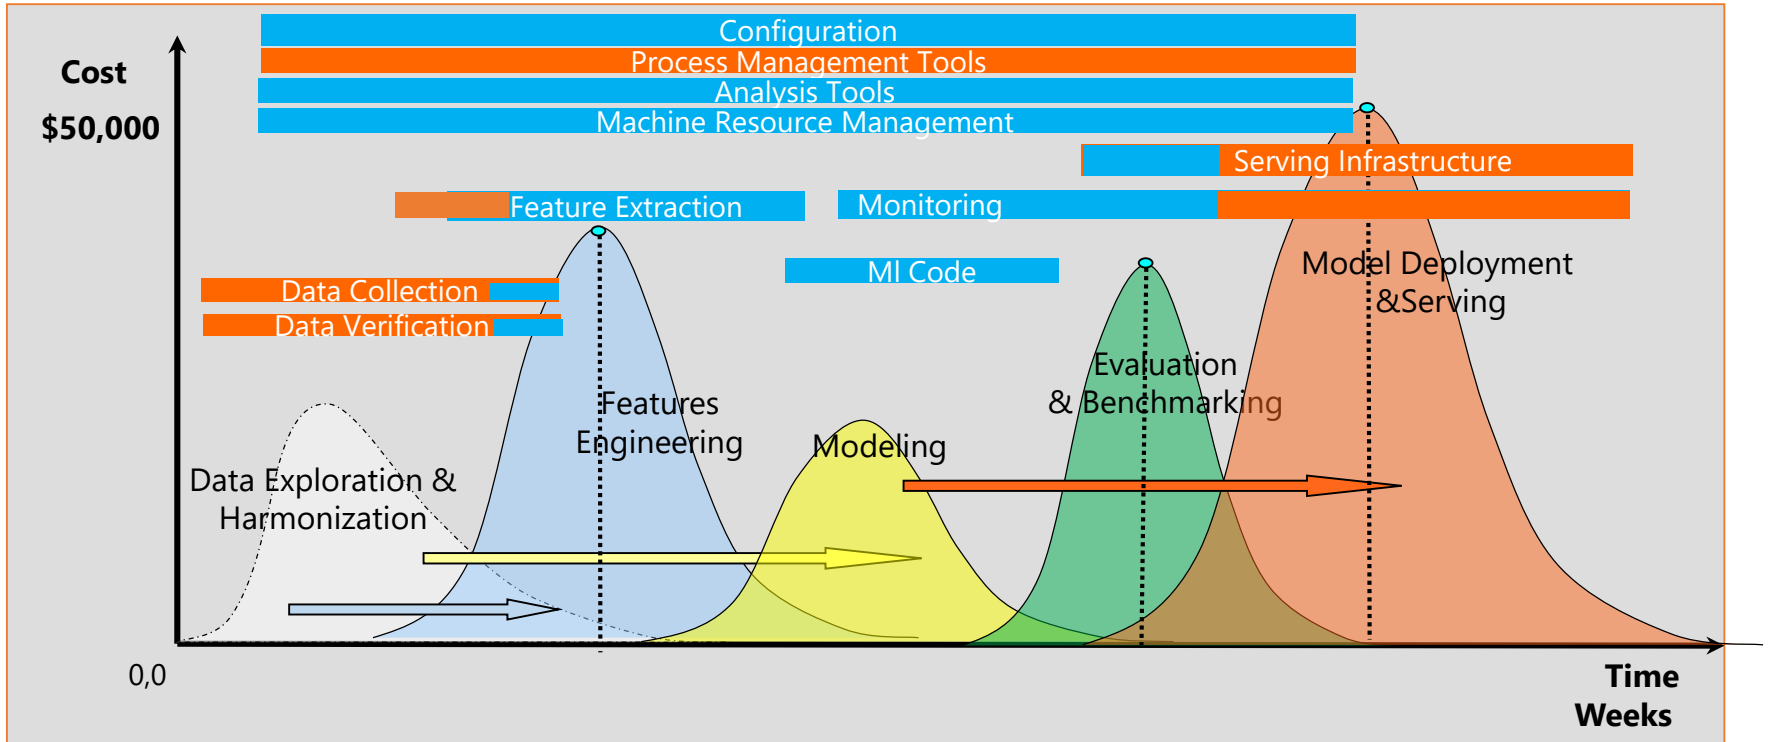
Journey 1: Machine learning frameworks

Example : Machine Learning Sandbox



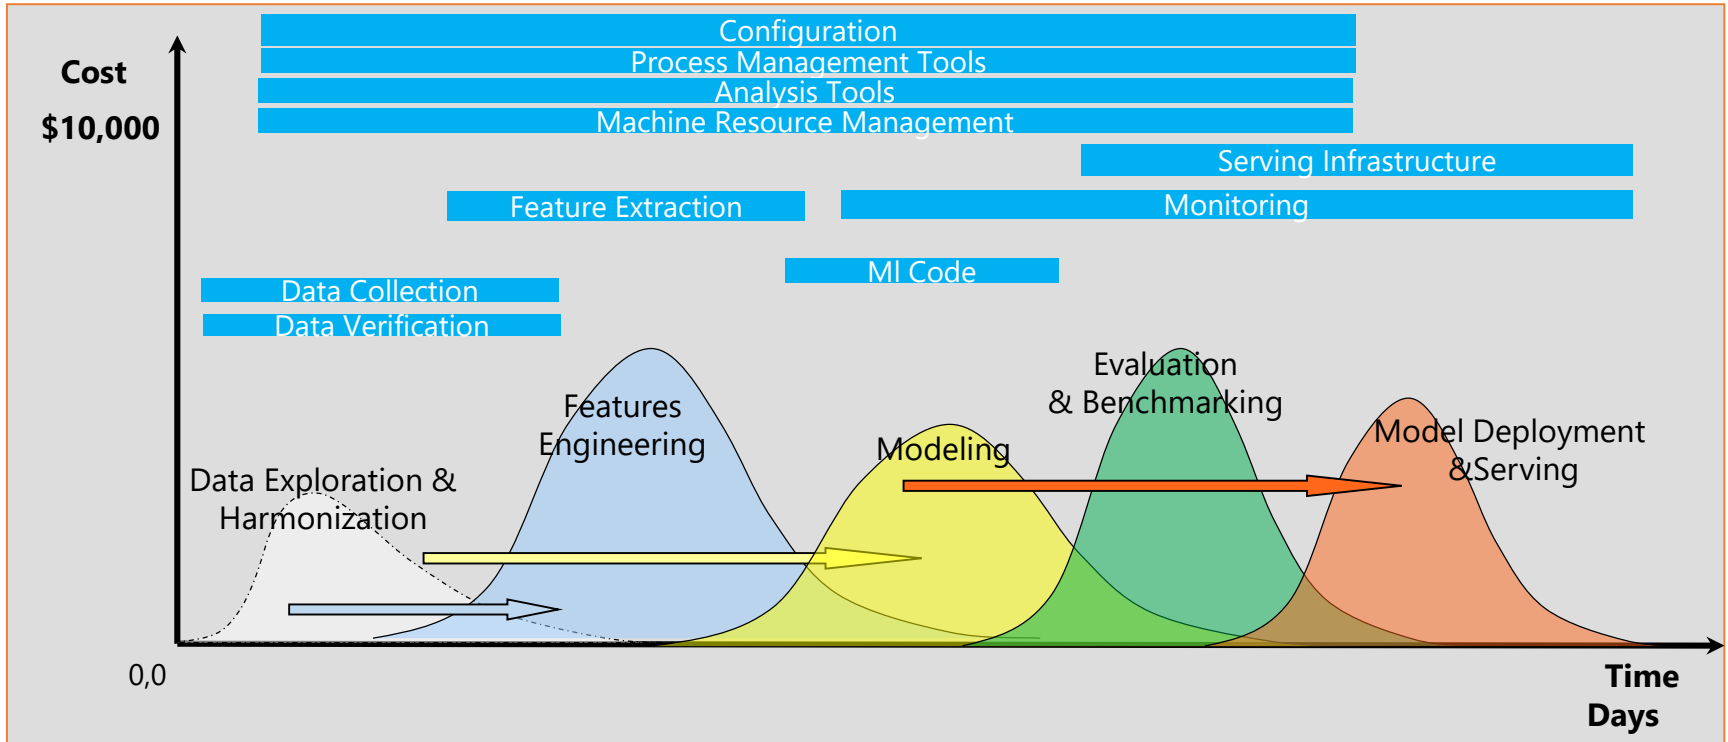
Journey 2: Fully managed machine learning services

Example : Data Science Workbench (Microsoft ML Studio)



Journey 3: Automation Services

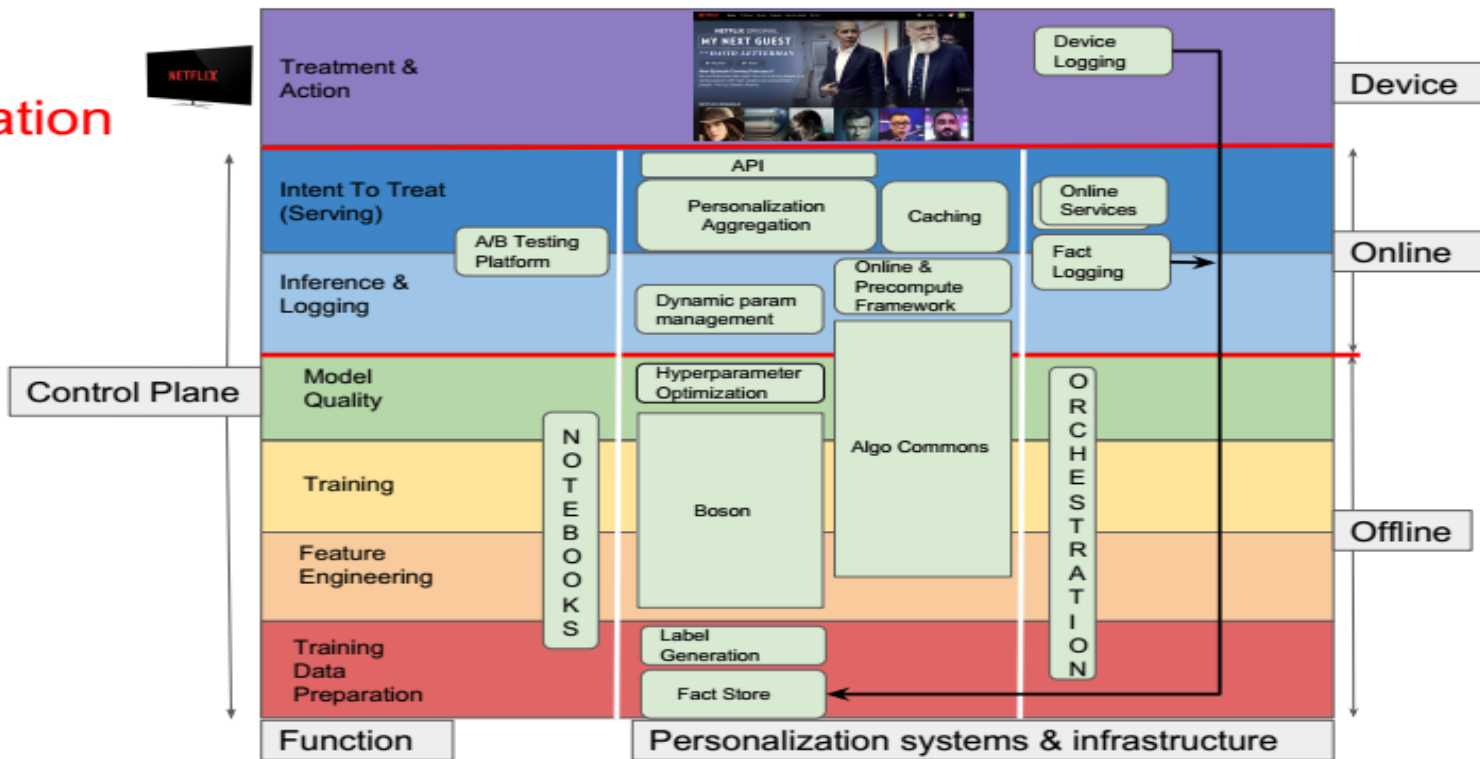
Example : Netflix , Amazon Sage Maker...



What we can learn from Netflix ?

<https://www.slideshare.net/FaisalZakariaSiddiqi/ml-infra-for-netflix-recommendations-ai-nextcon-talk>

The Personalization Rainbow



What we can learn from Amazon ?

Amazon SageMaker Components



Building

Training

Hosting

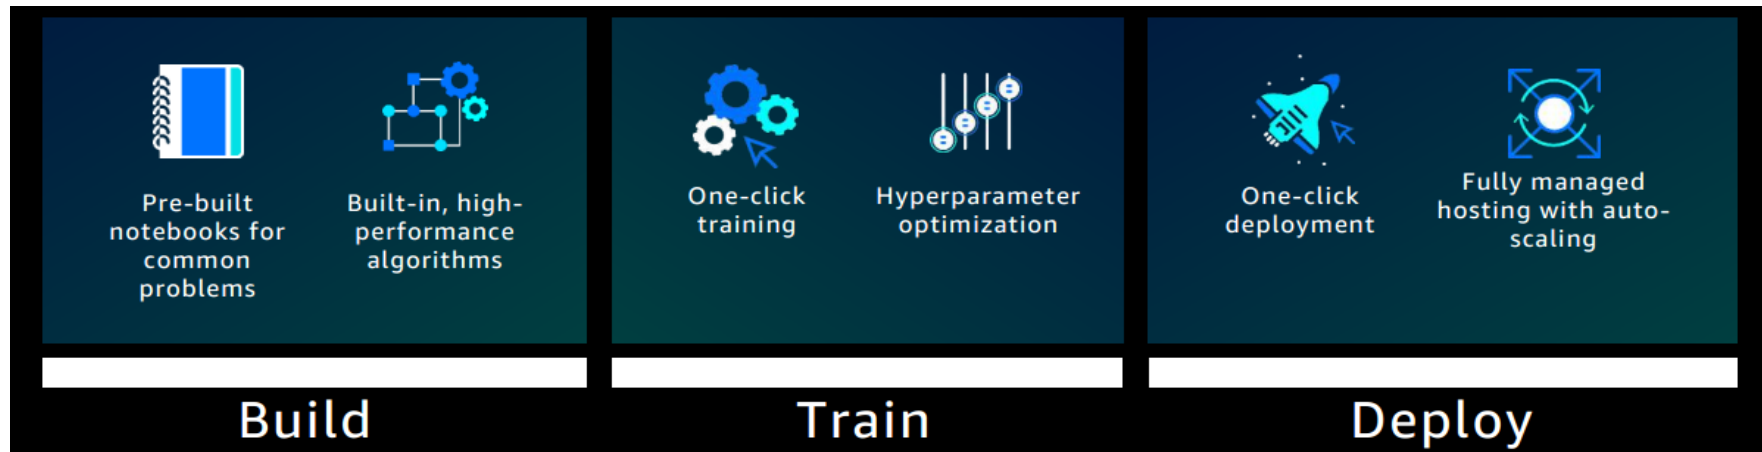
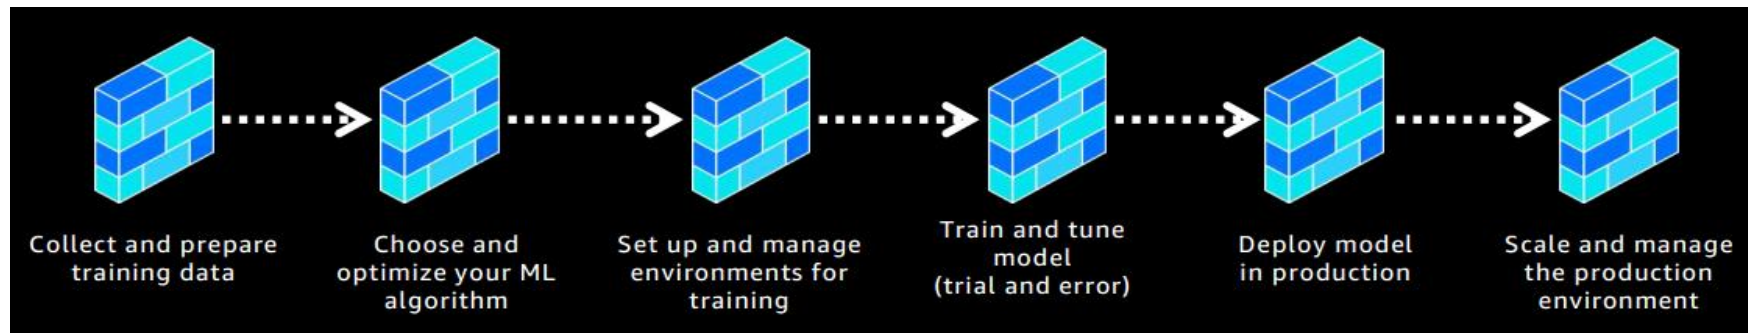
Amazon's fast, scalable algorithms

Distributed Apache MXNet and TensorFlow

Bring your own algorithm

Hyperparameter optimization

What we can learn from Amazon ? -- Continue



What we can learn from Prediction IO ?

<http://predictionio.apache.org/>

Apache PredictionIO® Documentation
Getting Started
Integrating with Your App
Deploying an Engine
Customizing an Engine
Collecting and Analyzing Data
Choosing an Algorithm(s)

Engine Template Gallery

 [Edit this page](#)

Pick a tab for the type of template you are looking for. Some still need to be ported (a simple process) to Apache PIO and these are marked. Also see each Template description for special support instructions.

Recommenders

Classification

Regression

NLP

Clustering

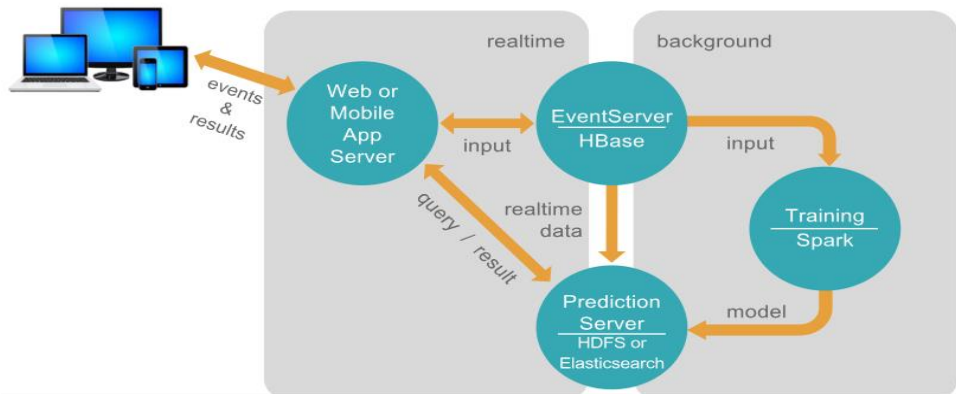
Similarity

Other

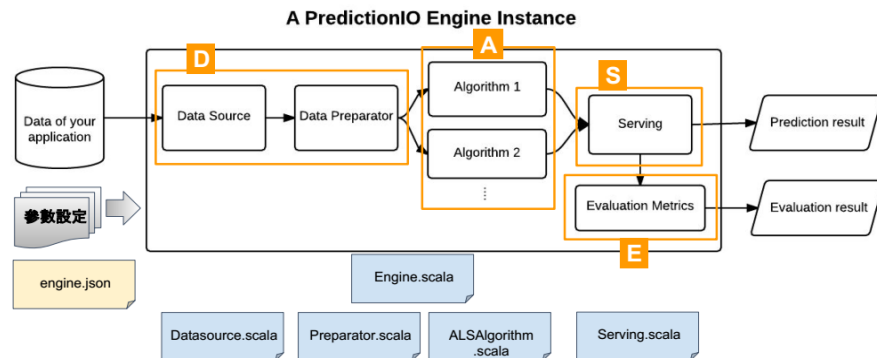
The Universal Recommender

 Star 404

Use for:



Customizing your Engine with D-A-S-E



Challenges of traditional machine learning

Learning

1

Feature engineering bottlenecks

Pre-calculate hundreds or thousands Long Term Variables take lots of resources and times (greater than 70%)

2

Model scalability limitations

Trade-off between automation in parallel and scaling machine learning to ever larger datasets and ever more complicated models

3

Heavily relies on human machine learning experts

Relies on human to perform the most of tasks, such as features selection, model selection, model hyper parameters tuning , critically analyze the results obtained...

Serving

4

Less integration with end to end data pipelines and serving multiple contexts

Gap to bring machine learning process into the existing enterprise data pipelines and application contexts including offline, streaming and real-time

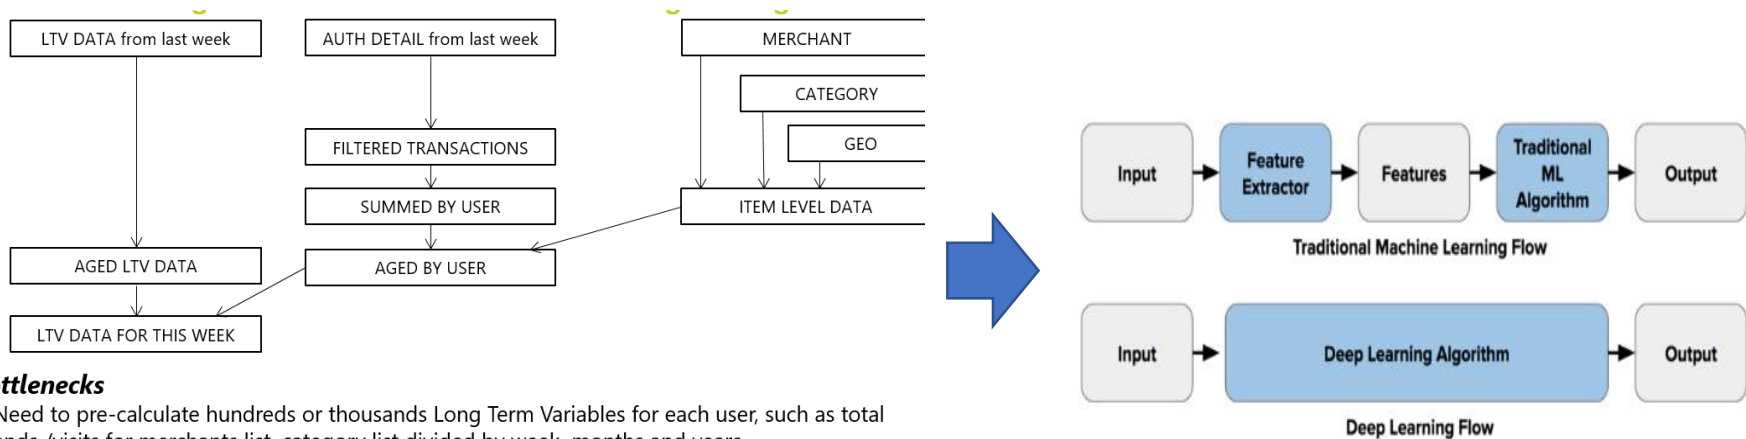
5

Isolated non-reusable APIs and CI/CD pipelines

Application teams spent lots of time with data scientist to build high level serving APIs , back and force. Most of time are using different technical stacks and non-reusable pipelines.

Isolated promotions and operation readiness without automate CI/CD

How deep learning can help -- Feature engineering bottlenecks




Bottlenecks

- Need to pre-calculate hundreds or thousands Long Term Variables for each user, such as total spends /visits for merchants list, category list divided by week, months and years
- The computation time for LTV features took > 70% of the data processing time for the whole lifecycle and occupied lots of resources which had huge impact to other critical workloads.
- Miss the feature selection optimizations which could save the data engineering efforts a lot

Improvements

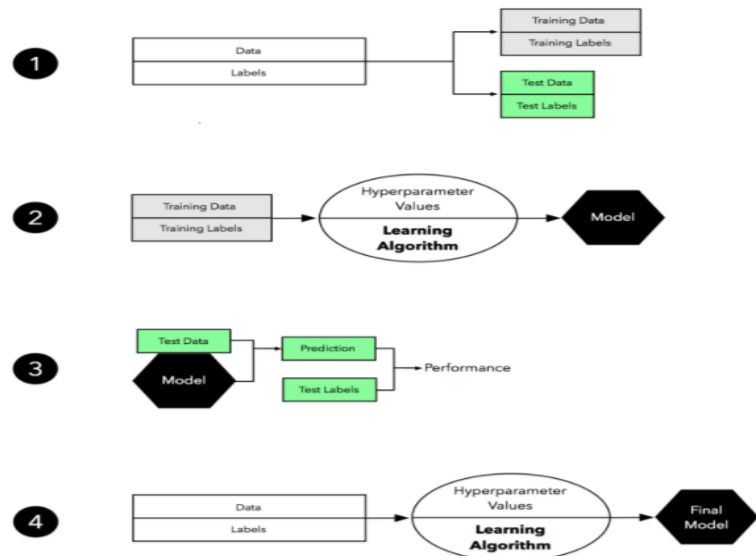
- When build model , only focus on few pre-defined sliding features and custom overlap features (Users only need to identify the columns names from data source)
- Remove most of the LTV pre-calculations works, saved hours time and lots of resources
- Deep learning algorithm generates exponential growth of hidden embedding features ,do the internal features selections and optimization automatically when it does cross validation at training stage

How deep learning can help -- Heavily relies on human machine learning experts



Relies on human to perform the following tasks:


- Select and construct appropriate features.
- Select an appropriate model family.
- Optimize model hyper parameters.
- Post process machine learning models.
- Critically analyze the results obtained.



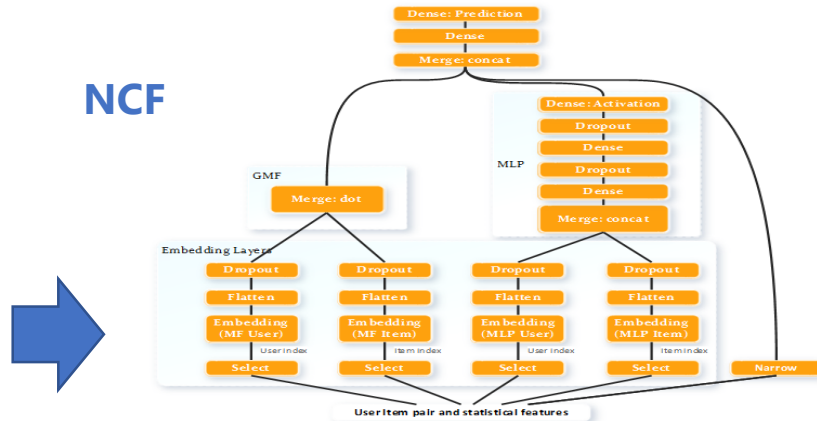
Improvements

- Common neural network "tricks", including initialization, L2 and dropout regularization, Batch normalization, gradient checking
- A variety of optimization algorithms, such as mini-batch gradient descent, Momentum, RMSprop and Adam
- Provides optimization-as-a-service using an ensemble of optimization strategies, allowing practitioners to efficiently optimize models faster and cheaper than standard approaches.


How deep learning can help -- Model scalability



NCF



Wide And Deep




- All the pipelines separated by items and generate one model for each item
- Have to pre-calculate the correlation matrix between items
- Lots of redundant duplications and computations at feature engineering, training and testing process
- Run items in parallel and occupied most of cluster resources when executed
- Bad metrics for items with few transactions
- It is very hard to scale more items, from hundreds to millions?

Improvements

- Scale models in deeper and wider without decreasing metrics

Enterprise requirements for AI as a Service




Collocated with mass data storage

- Analyze a large amount of data on the same Big Data clusters where the data are stored (HDFS, HBase, Hive, etc.) rather than **move or duplicate data**



Data governance with restricted Processing

- Follow data privacy, regulation and compliance (such as PCI/PII compliance and GDPR) rather than **operate data in unsecured zones**




Adopt structured sparse data sets

- More challenges also more benefits to adopt structured high dimensional sparse data sets rather than **non-structured dense data sets**



Seamless integration with Products Internal & External

- Add deep learning capabilities to existing Analytic Applications and/or machine learning workflows rather than **rebuild all of them**



Shared infrastructure with Multi-tenant isolated resources

- Leverage existing Big Data clusters and deep learning workloads should be managed and monitored with other workloads (ETL, data warehouse, traditional ML etc..) rather than **run ML/DL workloads standalone in separate clusters**



Automation and easy to go

- End to end automation rather than **manual efforts**
- Easy API enablement rather than **big learning curve or depend on special experts**

Deep learning approaches evaluation -- Super Stars



- Examples are good for dense high sample size data sets (But won't help us)
- Claimed that the GPU computing are better than CPU which requires new hardware infrastructure (very long timeline normally)
- Success requires many engineer-hours (Impossible to Install a Tensor Flow Cluster at STAGE ...)
- Low level APIs with steep learning curve (Where is your PhD degree ?)
- Not well integrated with other enterprise tools and need data movements (couldn't leverage the existing ETL, data warehousing and other analytic relevant data pipelines, technologies and tool sets. And it is also a big challenge to make duplicate data pipelines and data copy to the capacity and performance.)
- Tedious and fragile to distribute computations (less monitoring)
- The concerns of Enterprise Maturity and InfoSec (such as use GPU cluster with Tensor Flow from Google Cloud)

.....

Module 2

Keras introduction

-- Introduce a simple play ground , not just for Keras

<https://github.com/ufoym/deepo>



build passing docker pulls 49k license MIT

Deepo is a series of *Docker* images that

- allows you to quickly set up your deep learning research environment
- supports almost all *commonly used deep learning frameworks*
- supports *GPU acceleration* (CUDA and cuDNN included), also works in *CPU-only mode*
- works on Linux (*CPU version/GPU version*), Windows (*CPU version*) and OS X (*CPU version*)

and their Dockerfile generator that

- allows you to *customize your own environment* with Lego-like modules
- automatically resolves the dependencies for you

Keras introduction

-- Basic concepts

Keras: an API for specifying & training differentiable programs

Keras API

TensorFlow / CNTK / MXNet / Theano / ...

GPU

CPU

TPU

Keras is the official high-level API of TensorFlow

tf.keras

TensorFlow

GPU

CPU

TPU

Keras introduction

-- Three API styles

The Sequential Model

- - Dead simple
- - Only for single-input, single-output, sequential layer stacks
- - Good for 70+% of use cases

The functional API

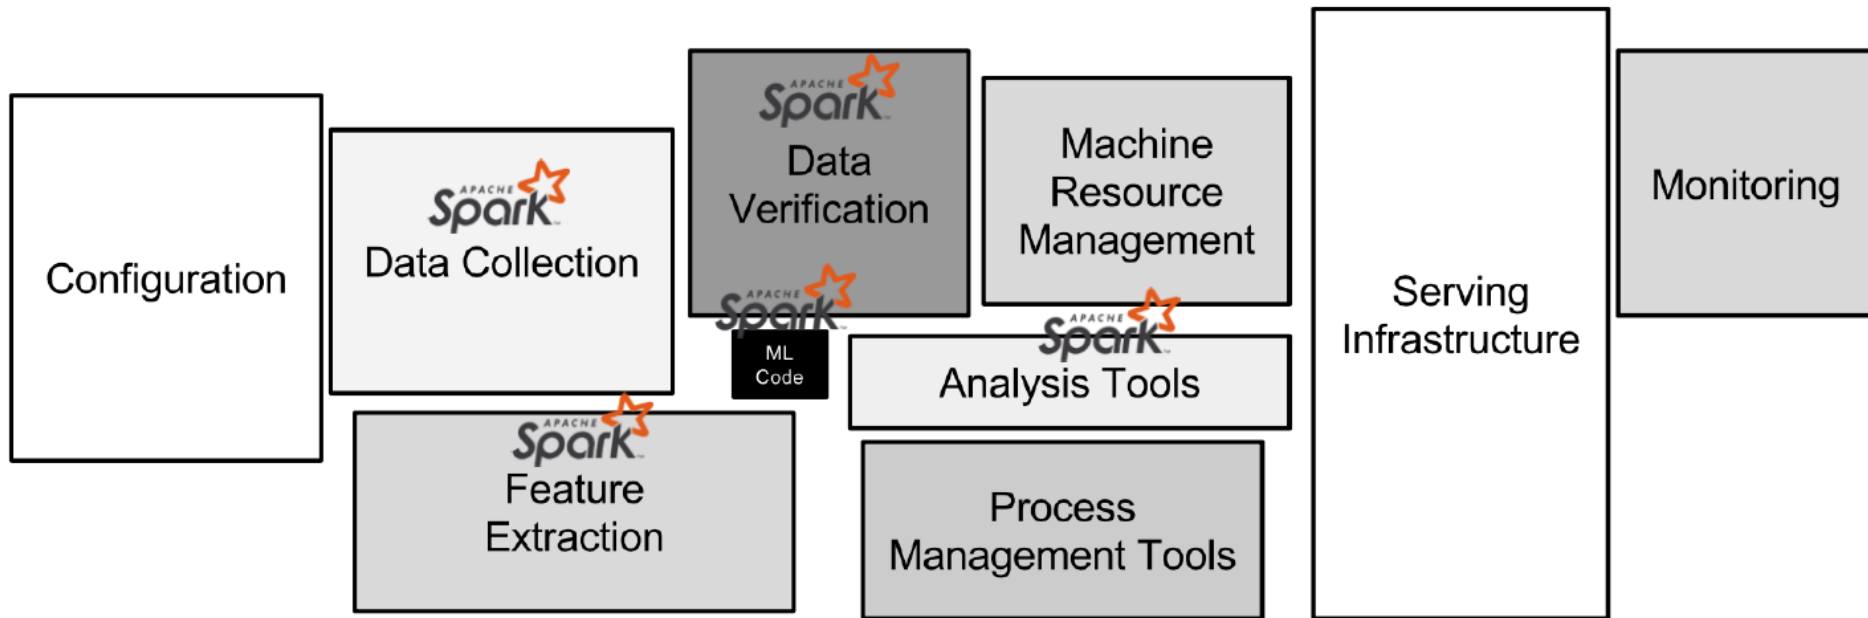
- - Like playing with Lego bricks
- - Multi-input, multi-output, arbitrary static graph topologies
- - Good for 95% of use cases

Model subclassing

- - Maximum flexibility
- - Larger potential error surface

Code time !

Why Spark ?



What does Spark offer for deep learning ?

Integrations with existing DL libraries

- Deep Learning Pipelines (from Databricks)
- Caffe (CaffeOnSpark)
- Keras (Elephas)
- mxnet
- Paddle
- TensorFlow (TensorFlow on Spark, TensorFrames)
- CNTK (mmlspark)

Implementations of DL on Spark

- BigDL+ Analytic Zoo
- DeepDist
- DeepLearning4J
- SparkCL
- SparkNet

Options of Keras on Spark

 maxpumperla / elephas

 Watch ▾

108

★ Star

1,122

 Fork

246

↔ Code

! Issues 32

 Pull requests 0

 Projects 0

 Wiki

 Insights

Distributed Deep learning with Keras & Spark <http://maxpumperla.com/elephas/>


spark

keras

neural-networks

deep-learning

distributed-computing

 309 commits


 3 branches

 5 releases

 16 contributors

 MIT

 intel-analytics / BigDL

 Watch ▾

220

★ Star

2,776

 Fork

662


↔ Code

! Issues 109

 Pull requests 30

 Projects 0

 Wiki

 Insights

BigDL: Distributed Deep Learning Library for Apache Spark <https://bigdl-project.github.io/>

deep-learning

spark

neural-network

big-data

hadoop

python

scala

keras


ai

 2,442 commits

 11 branches

 8 releases

 54 contributors

 Apache-2.0

Why BigDL ?

Lenet	Inception	Vgg	Resnet
Fast RCNN	SSD	RNN	LSTM
Deep Speech	Seq2Seq	Auto Encoder	Recommendation

Examples	Documents
Notebook	Tensor Board
Scala	Python

Layers

Spatial Convolution	Spatial MaxPooling	Volumetric Convolution	Volumetric MaxPooling	Fully Connected
ReLu	LRN	RNN	Batch Normalization	Other Layers

Optimization

SGD	Adagrad	LBFGS	Adamx
Adam	Adadelta	RMSprop	

Criterion

Cross Entropy	ClassNLL	MSE
Dice Coefficient	Margin	Abs

Deep Learning Building Blocks
Layers, Optimizers, Criterion
Deep Learning Models

Scala* and Python* support
Spark ML Pipeline integration
Jupyter* notebook integration
Tensorboard* integration
OpenCV* support
Model Interoperability
(Caffe*/TensorFlow*/Keras*)



CONSUMER

CALL CENTER ROUTING
IMAGE SIMILARITY SEARCH
SMART JOB SEARCH



HEALTH

ANALYSIS OF 3D MRI
MODELS FOR KNEE
DEGRADATION



FINANCE

FRAUD DETECTION
RECOMMENDATION
CUSTOMER/MERCHANT
PROPENSITY



RETAIL

IMAGE FEATURE
EXTRACTION



MANUFACTURING

STEEL SURFACE
DEFECT DETECTION



SCIENTIFIC COMPUTING

WEATHER
FORECASTING

Why Analytics Zoo ?

Analytics Zoo -> Unified Analytics + AI Platform for Spark and BigDL

Reference Use Cases	Anomaly detection, sentiment analysis, fraud detection, chatbot, sequence prediction, etc.
Built-In Algorithms and Models	Image classification, object detection, text classification, recommendations, GAN, etc.
Feature Engineering and Transformations	Image, text, speech, 3D imaging, time series, etc.
High-Level Pipeline APIs	DataFrames, ML Pipelines, Autograd, Transfer Learning, etc.
Runtime Environment	Spark, BigDL, Python, etc.

<https://github.com/intel-analytics/analytics-zoo>

Use case for user item propensity model

- Propensity to buy
- Propensity to use
- Propensity to engage
- Propensity to contract
- Etc.

- Life Insurance
- Auto Insurance
- Homeowner's Insurance
- Mortgage
- Re-Financing
- Credit cards
- Personal Loans
- Investments
- Satellite TV
- Cable
- Netflix
- Online banking
- Online money management
- Voting
- Disease



<https://catalog.data.gov/dataset/purchase-card-pcard-fiscal-year-2014>


Purchase Card (PCard) Fiscal Year 2014

□ **Metadata Updated:** September 15, 2016

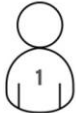


This dataset contains information on purchases made through the purchase card programs administered by the state and higher ed institutions. The purchase card information will be updated monthly after the end of the month. For example, July information will be added in August.

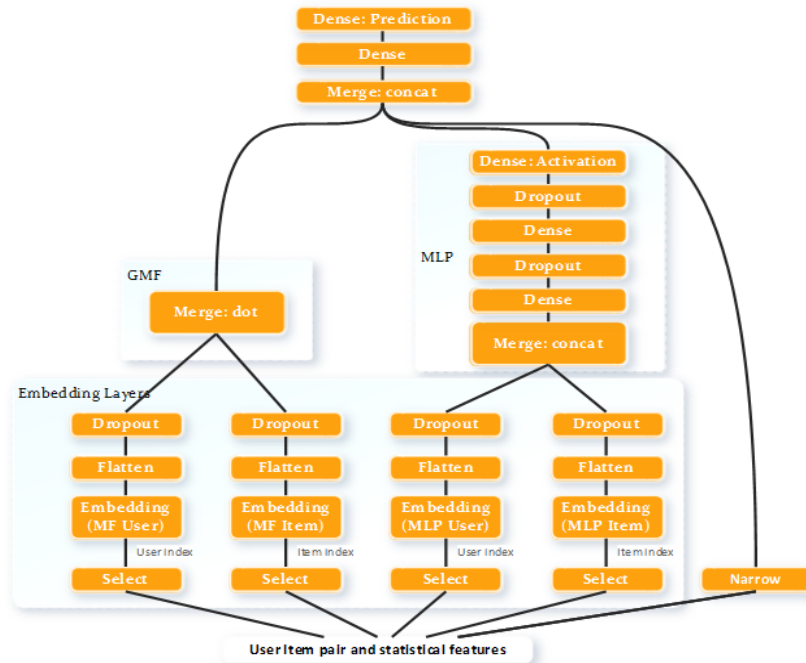
Build a User Item Propensity model with deep learning algorithms

Logistic regression



Collaborative Filtering

			
Product A	5★	5★	5★
Product B	3★	3★	3★
Product C	5★	3★	4★



Neural Collaborative Filtering deep learning algorithm

<https://arxiv.org/pdf/1708.05031.pdf>

https://github.com/hexiangnan/neural_collaborative_filtering

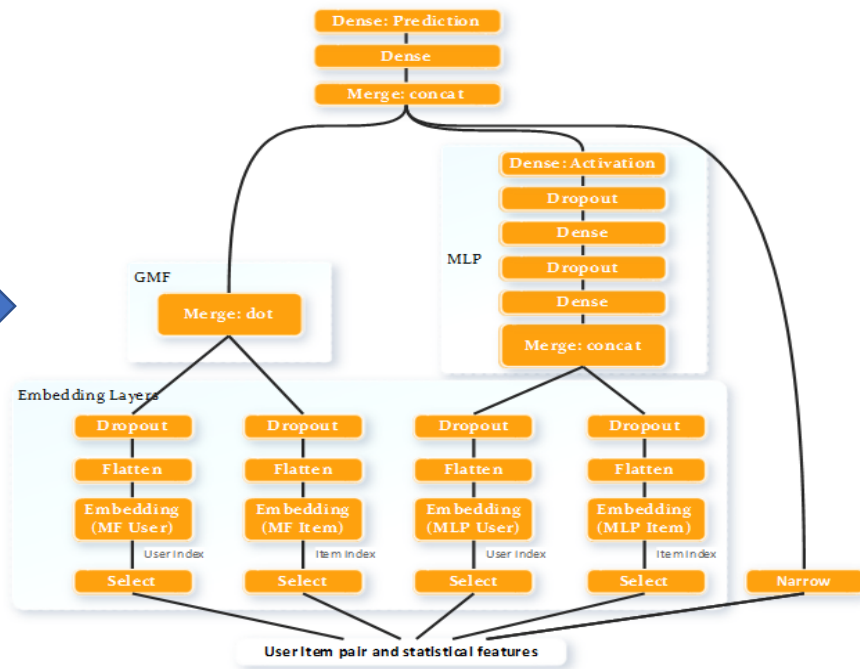
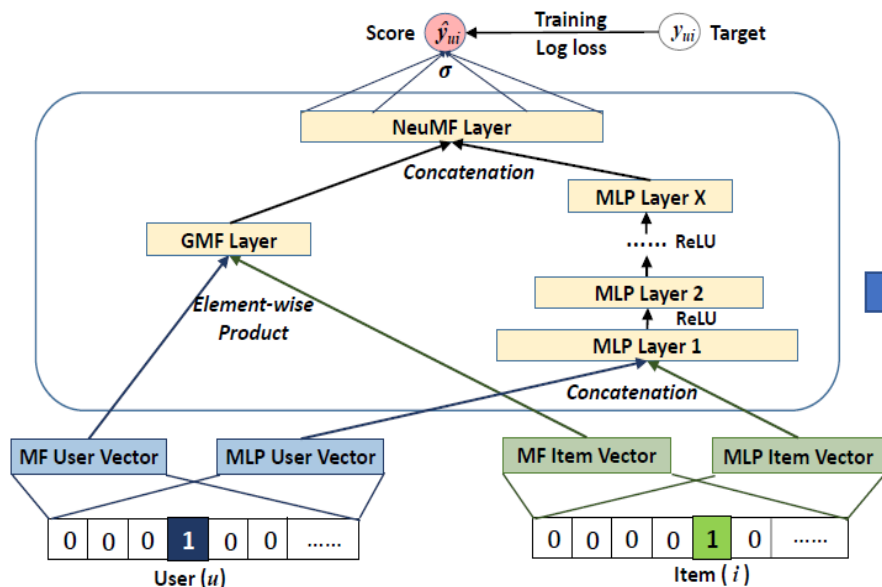


Figure 3: Neural matrix factorization model

Code Lab 1

Build a Docker image and run a Keras on Spark container

<https://github.com/jack1981/AaaS Demo>

Run the NCF deep learning pipeline for User Item Propensity model

<https://github.com/jack1981/AaaSDemo>

Q & A